

# Classification

Name: Ryan Donaldson

Date: 02/13/2023

## Summary

Linear models predict which category a given input belongs to by creating a linear decision to separate different data classes from each other. The strength of linear models is that they can be used with large datasets and are easy to visualize and implement. The weakness of linear models is that they assume a linear relationship between the features and the target variable.

Please click here (<https://www.kaggle.com/datasets/yasserh/uber-fares-dataset>) to access the dataset used in this project.

## Data Cleaning

Before exploring and visualizing data, let's clean our dataset to make our overall linear regression model easier to work with as well as compiling informative graphs later on.

First, we will read the CSV and load the geosphere package so we can better work with latitude and longitude data that the columns will provide. We will also load the dplyr package to make column and row operations and mutations easier. We will load the e1071 package to build a Naive Bayes model and caret for useful distribution statistics.

```
Uber <- read.csv("uber.csv", na.strings="NA", header=TRUE)
if(!require("geosphere")) {
  install.packages("geosphere")
  library("geosphere")
}
```

```
## Loading required package: geosphere
```

```
if(!require("dplyr")) {
  install.packages("dplyr")
  library("dplyr", warn.conflicts=FALSE)
}
```

```
## Loading required package: dplyr
```

```
##
## Attaching package: 'dplyr'
```

```
## The following objects are masked from 'package:stats':  
##  
##   filter, lag
```

```
## The following objects are masked from 'package:base':  
##  
##   intersect, setdiff, setequal, union
```

```
if(!require("e1071")) {  
  install.packages("e1071")  
  library("e1071", warn.conflicts=FALSE)  
}
```

```
## Loading required package: e1071
```

```
if(!require("caret")) {  
  install.packages("caret")  
  library("caret", warn.conflicts=FALSE)  
}
```

```
## Loading required package: caret
```

```
## Loading required package: ggplot2
```

```
## Loading required package: lattice
```

Let's remove any rows which contain any latitude or longitude outliers.

```

exclude_pickup_lon <- filter(Uber, pickup_longitude < -360
  | pickup_longitude > 360
  | pickup_longitude < -180
  | pickup_longitude > 180)
Uber <- anti_join(Uber, exclude_pickup_lon, by="key")

exclude_dropoff_lon <- filter(Uber, dropoff_longitude < -360
  | dropoff_longitude > 360
  | dropoff_longitude < -180
  | dropoff_longitude > 180)
Uber <- anti_join(Uber, exclude_dropoff_lon, by="key")

exclude_pickup_lat <- filter(Uber, pickup_latitude < -90
  | pickup_latitude > 90)
Uber <- anti_join(Uber, exclude_pickup_lat, by="key")

exclude_dropoff_lat <- filter(Uber, dropoff_latitude < -90
  | dropoff_latitude > 90)
Uber <- anti_join(Uber, exclude_dropoff_lat, by="key")

```

Now, let's mutate the data frame and add a new Distance column. Each row will have a value representing the shortest distance between its corresponding pickup and dropoff points. We will use the `distHaversine` function from `geosphere` for the math behind this. The `Distance_In_Km` will be represented in kilometers from the function output.

```
print("Adding Distance column, please wait...")
```

```
## [1] "Adding Distance column, please wait..."
```

```

Uber <- Uber %>% rowwise() %>%
  mutate(Distance_In_Km=distHaversine(c(pickup_longitude, pickup_latitude), c(dropoff_longitude, dropoff_latitude)))
print("Uber dataframe has been mutated")

```

```
## [1] "Uber dataframe has been mutated"
```

Now, we will divide the data into train and test sets using a 80/20 split.

```

set.seed(1234)
i <- sample(1:nrow(Uber), nrow(Uber)*0.80,
  replace=FALSE)
train <- Uber[i,]
test <- Uber[-i,]

```

## Data Exploration

Next, we will run 5 R functions for data exploration of the data set using the training data. First, let's run the `str()` function to get a look into the format of the data.

```
str(train)
```

```

## rows_df [159,990 × 10] (S3: rowwise_df/tbl_df/tbl/data.frame)
## $ X : int [1:159990] 20105724 18896163 40414748 11607434 34954013 290
82183 27715186 29305241 29942390 22199820 ...
## $ key : chr [1:159990] "2010-03-04 08:53:30.0000006" "2014-10-12 03:30:
41.0000002" "2010-07-16 07:57:00.00000030" "2013-09-14 09:27:00.000000131" ...
## $ fare_amount : num [1:159990] 9.7 17 34.3 7.5 13.5 ...
## $ pickup_datetime : chr [1:159990] "2010-03-04 08:53:30 UTC" "2014-10-12 03:30:41 U
TC" "2010-07-16 07:57:00 UTC" "2013-09-14 09:27:00 UTC" ...
## $ pickup_longitude : num [1:159990] -73.8 -74 -73.9 -74 -74 ...
## $ pickup_latitude : num [1:159990] 40.8 40.7 40.8 40.7 40.8 ...
## $ dropoff_longitude: num [1:159990] -73.8 -73.9 -74 -74 -74 ...
## $ dropoff_latitude : num [1:159990] 40.8 40.8 40.8 40.7 40.8 ...
## $ passenger_count : int [1:159990] 1 1 1 1 1 3 1 1 1 1 ...
## $ Distance_In_Km : num [1:159990] 3048 6429 9321 1853 5147 ...
## - attr(*, "groups")= tibble [159,990 × 1] (S3: tbl_df/tbl/data.frame)
## ..$ .rows: list<int> [1:159990]
## .. ..$ : int 1
## .. ..$ : int 2
## .. ..$ : int 3
## .. ..$ : int 4
## .. ..$ : int 5
## .. ..$ : int 6
## .. ..$ : int 7
## .. ..$ : int 8
## .. ..$ : int 9
## .. ..$ : int 10
## .. ..$ : int 11
## .. ..$ : int 12
## .. ..$ : int 13
## .. ..$ : int 14
## .. ..$ : int 15
## .. ..$ : int 16
## .. ..$ : int 17
## .. ..$ : int 18
## .. ..$ : int 19
## .. ..$ : int 20
## .. ..$ : int 21
## .. ..$ : int 22
## .. ..$ : int 23
## .. ..$ : int 24
## .. ..$ : int 25
## .. ..$ : int 26
## .. ..$ : int 27
## .. ..$ : int 28
## .. ..$ : int 29
## .. ..$ : int 30
## .. ..$ : int 31
## .. ..$ : int 32
## .. ..$ : int 33
## .. ..$ : int 34
## .. ..$ : int 35
## .. ..$ : int 36

```

```
## .. ..$ : int 37
## .. ..$ : int 38
## .. ..$ : int 39
## .. ..$ : int 40
## .. ..$ : int 41
## .. ..$ : int 42
## .. ..$ : int 43
## .. ..$ : int 44
## .. ..$ : int 45
## .. ..$ : int 46
## .. ..$ : int 47
## .. ..$ : int 48
## .. ..$ : int 49
## .. ..$ : int 50
## .. ..$ : int 51
## .. ..$ : int 52
## .. ..$ : int 53
## .. ..$ : int 54
## .. ..$ : int 55
## .. ..$ : int 56
## .. ..$ : int 57
## .. ..$ : int 58
## .. ..$ : int 59
## .. ..$ : int 60
## .. ..$ : int 61
## .. ..$ : int 62
## .. ..$ : int 63
## .. ..$ : int 64
## .. ..$ : int 65
## .. ..$ : int 66
## .. ..$ : int 67
## .. ..$ : int 68
## .. ..$ : int 69
## .. ..$ : int 70
## .. ..$ : int 71
## .. ..$ : int 72
## .. ..$ : int 73
## .. ..$ : int 74
## .. ..$ : int 75
## .. ..$ : int 76
## .. ..$ : int 77
## .. ..$ : int 78
## .. ..$ : int 79
## .. ..$ : int 80
## .. ..$ : int 81
## .. ..$ : int 82
## .. ..$ : int 83
## .. ..$ : int 84
## .. ..$ : int 85
## .. ..$ : int 86
## .. ..$ : int 87
## .. ..$ : int 88
```

```
## .. ..$ : int 89
## .. ..$ : int 90
## .. ..$ : int 91
## .. ..$ : int 92
## .. ..$ : int 93
## .. ..$ : int 94
## .. ..$ : int 95
## .. ..$ : int 96
## .. ..$ : int 97
## .. ..$ : int 98
## .. ..$ : int 99
## .. .. [list output truncated]
## .. ..@ ptype: int(0)
```

Next, let's gather an overall basic summary of each column of our training data.

```
summary(train)
```

```
##           X           key      fare_amount      pickup_datetime
## Min.      :      1      Length:159990      Min.      : -52.00      Length:159990
## 1st Qu.:13886062      Class :character      1st Qu.:   6.00      Class :character
## Median :27763618      Mode  :character      Median :   8.50      Mode  :character
## Mean     :27749116                                Mean     : 11.37
## 3rd Qu.:41600315                                3rd Qu.: 12.50
## Max.     :55423567                                Max.     :499.00
##
## pickup_longitude pickup_latitude dropoff_longitude dropoff_latitude
## Min.      : -93.82      Min.      : -74.01      Min.      : -75.42      Min.      : -74.02
## 1st Qu.: -73.99      1st Qu.:  40.73      1st Qu.: -73.99      1st Qu.:  40.73
## Median : -73.98      Median :  40.75      Median : -73.98      Median :  40.75
## Mean     : -72.51      Mean     :  39.93      Mean     : -72.52      Mean     :  39.93
## 3rd Qu.: -73.97      3rd Qu.:  40.77      3rd Qu.: -73.96      3rd Qu.:  40.77
## Max.     :  40.81      Max.     :  48.02      Max.     :  40.83      Max.     :  45.03
##                                     NA's      :1          NA's      :1
## passenger_count      Distance_In_Km
## Min.      :   0.000      Min.      :      0
## 1st Qu.:   1.000      1st Qu.:   1216
## Median :   1.000      Median :   2127
## Mean     :   1.683      Mean     :  20083
## 3rd Qu.:   2.000      3rd Qu.:   3886
## Max.     : 208.000      Max.     :8792737
##                                     NA's      :1
```

Our summary above tells us that exploring missing data is really not necessary considering the `dropoff_longitude` and `dropoff_latitude` columns each only contain 1 NA. So, first let's just look at the first few rows.

```
head(train)
```

<b>X</b>	<b>key</b>	<b>fare_amount</b>	<b>pickup_datetime</b>	<b>pickup_longitude</b>	<b>pickup_latitude</b>	<b>passenger_count</b>
<int>	<chr>	<dbl>	<chr>	<dbl>	<dbl>	<int>
20105724	2010-03-04 08:53:30.0000006	9.70	2010-03-04 08:53:30 UTC	-122.509392	37.80901	1
18896163	2014-10-12 03:30:41.0000002	17.00	2014-10-12 03:30:41 UTC	-122.509392	37.80901	1
40414748	2010-07-16 07:57:00.00000030	34.27	2010-07-16 07:57:00 UTC	-122.509392	37.80901	1
11607434	2013-09-14 09:27:00.000000131	7.50	2013-09-14 09:27:00 UTC	-122.509392	37.80901	1
34954013	2014-06-20 06:19:25.0000003	13.50	2014-06-20 06:19:25 UTC	-122.509392	37.80901	1
29082183	2011-03-27 17:09:00.000000162	49.80	2011-03-27 17:09:00 UTC	-122.509392	37.80901	1

6 rows | 1-5 of 10 columns

Now, let's look at the last few rows of our training data.

```
tail(train)
```

<b>X</b>	<b>key</b>	<b>fare_amount</b>	<b>pickup_datetime</b>	<b>pickup_longitude</b>	<b>pickup_latitude</b>	<b>passenger_count</b>
<int>	<chr>	<dbl>	<chr>	<dbl>	<dbl>	<int>
55251256	2014-08-20 23:51:55.0000002	32.83	2014-08-20 23:51:55 UTC	-122.509392	37.80901	1
9135184	2012-07-09 08:28:31.0000001	4.10	2012-07-09 08:28:31 UTC	-122.509392	37.80901	1
38983892	2010-06-25 12:27:48.0000004	4.50	2010-06-25 12:27:48 UTC	-122.509392	37.80901	1
17122458	2011-04-26 01:03:03.0000001	8.50	2011-04-26 01:03:03 UTC	-122.509392	37.80901	1
28892833	2012-06-06 15:12:44.0000001	31.30	2012-06-06 15:12:44 UTC	-122.509392	37.80901	1
37672342	2012-05-09 21:34:00.000000118	9.30	2012-05-09 21:34:00 UTC	-122.509392	37.80901	1

6 rows | 1-5 of 10 columns

Since we're dealing with over 200,000 records across 9 columns. Let's also explore the column names within our data set in case we need to reference them later when making predictions.

```
names(train)
```

```
## [1] "X" "key" "fare_amount"
## [4] "pickup_datetime" "pickup_longitude" "pickup_latitude"
## [7] "dropoff_longitude" "dropoff_latitude" "passenger_count"
## [10] "Distance_In_Km"
```


## Data Visualization

Let's create some informative graphs based on this training data, particularly distributions of various columns since we're working with such a large dataset. First, let's get a sense of the distribution of ride distances.



```
boxplot(train$Distance_In_Km, col="slategray", horizontal=TRUE, xlab="Distance", main="Distance of Rides")
```

## Distance of Rides




Next,

let's get a sense of the distribution of fare amounts.

```
boxplot(train$fare_amount, col="slategray", horizontal=TRUE, xlab="Fare", main="Fare Amounts")
```

## Fare Amounts




Next,

let's get a sense of the distribution of passenger counts.

```
boxplot(train$passenger_count, col="slategray", horizontal=TRUE, xlab="Number of Passengers", main="Passengers Per Ride")
```

## Passengers Per Ride



looking at the distribution among these columns, we've identified some outliers that would be best to remove.

We cannot have negative fare amounts, zero passengers, more than 6 passengers and more than 0 km traveled in order to accurately represent a ride. We also should remove large distances as well (let's say more than 50 miles) as well as large fare amounts (let's say more than 60 dollars).

Let's further filter our Uber data set and then do a 80/20 split again.

```
exclude_outliers <- filter(Uber, fare_amount < 0
  | fare_amount == 0
  | fare_amount > 60
  | passenger_count == 0
  | passenger_count > 6
  | Distance_In_Km == 0
  | Distance_In_Km > 80)
Uber <- anti_join(Uber, exclude_outliers, by="key")
```

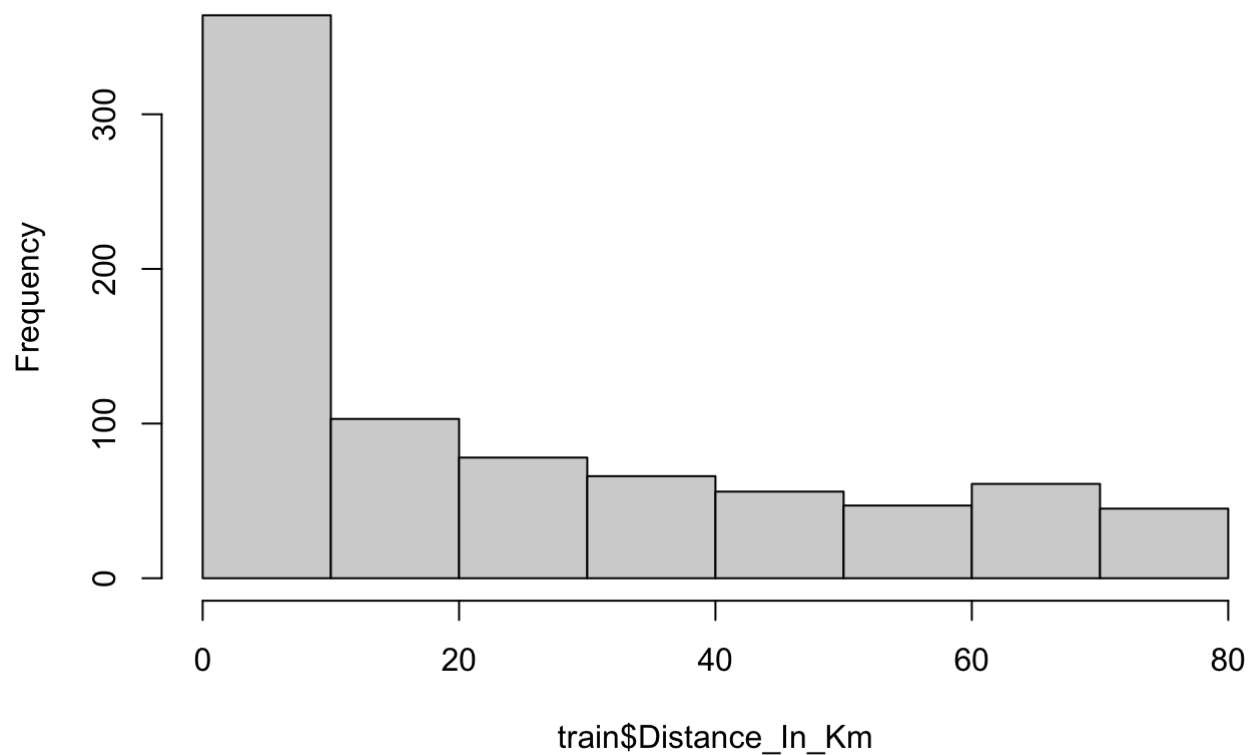
Now, let's divide the new filtered data into train and test sets using a 80/20 split.

```
set.seed(1234)
i <- sample(1:nrow(Uber), nrow(Uber)*0.80,
  replace=FALSE)
train <- Uber[i,]
test <- Uber[-i,]
```

Now that we've removed outliers, let's see a histogram of the distance of rides.

```
hist(train$Distance_In_Km)
```

### Histogram of train\$Distance\_In\_Km




Now

let's see a histogram corresponding to fare amounts.

```
hist(train$fare_amount)
```

## Histogram of train\$fare\_amount



###

Logistic Regression Next, we will build a logistic regression model and output the summary.

```
glm1 <- glm(as.factor(fare_amount)~Distance_In_Km, data=train, family=binomial)
summary(glm1)
```

```
##
## Call:
## glm(formula = as.factor(fare_amount) ~ Distance_In_Km, family = binomial,
##      data = train)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -3.4480    0.0024    0.0249    0.0708    0.0899
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)      5.4955      1.2420   4.425 9.66e-06 ***
## Distance_In_Km    0.1751      0.2880   0.608   0.543
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 15.417  on 819  degrees of freedom
## Residual deviance: 13.886  on 818  degrees of freedom
## AIC: 17.886
##
## Number of Fisher Scoring iterations: 13
```

First, the median residual indicates that most residuals are close to zero. Second, the AIC measures the overall quality of the model where lower values indicate that the model is a good fit. We can see that the AIC has a low value of 17.886. With these two factors in mind, we can hypothesize that logistic regression might be a good fit for the data.

## Naive Bayes

Next, we will build a Naive Bayes model and output what the model learned.

```
nb1 <- naiveBayes(fare_amount~Distance_In_Km, data=train)
nb1
```

```
##
## Naive Bayes Classifier for Discrete Predictors
##
## Call:
## naiveBayes.default(x = X, y = Y, laplace = laplace)
##
## A-priori probabilities:
## Y
##      0.01      0.11      2.5      2.9      3      3.25
## 0.001219512 0.001219512 0.365853659 0.025609756 0.042682927 0.001219512
##      3.3      3.4      3.5      3.7      3.8      4
## 0.017073171 0.001219512 0.018292683 0.013414634 0.001219512 0.004878049
##      4.1      4.5      4.9      5      5.1      5.3
## 0.010975610 0.036585366 0.007317073 0.007317073 0.001219512 0.007317073
##      5.5      5.7      6      6.1      6.2      6.3
## 0.003658537 0.008536585 0.002439024 0.002439024 0.001219512 0.001219512
##      6.5      6.9      7      7.3      7.5      7.7
## 0.020731707 0.006097561 0.004878049 0.003658537 0.003658537 0.006097561
##      8      8.1      8.2      8.5      8.6      8.9
## 0.004878049 0.008536585 0.001219512 0.008536585 0.001219512 0.004878049
##      9      9.3      9.5      9.6      9.7      9.8
## 0.004878049 0.003658537 0.004878049 0.001219512 0.002439024 0.001219512
##      10      10.1      10.2      10.3      10.5      10.9
## 0.004878049 0.003658537 0.001219512 0.001219512 0.007317073 0.001219512
##      11      11.3      11.5      11.7      12      12.1
## 0.003658537 0.007317073 0.001219512 0.002439024 0.003658537 0.007317073
##      12.5      12.9      13      13.3      13.5      13.8
## 0.003658537 0.002439024 0.002439024 0.001219512 0.003658537 0.001219512
##      14      14.1      14.5      14.6      14.9      15
## 0.003658537 0.003658537 0.004878049 0.001219512 0.003658537 0.004878049
##      15.5      15.7      15.9      16      16.5      16.9
## 0.003658537 0.001219512 0.001219512 0.001219512 0.004878049 0.002439024
##      17      17.3      17.5      17.9      18      18.1
## 0.003658537 0.001219512 0.007317073 0.001219512 0.001219512 0.001219512
##      18.9      19      19.3      19.7      19.8      20
## 0.001219512 0.001219512 0.001219512 0.001219512 0.001219512 0.015853659
##      20.5      20.9      23      23.7      24.5      25
## 0.002439024 0.002439024 0.001219512 0.001219512 0.001219512 0.002439024
##      27      28      28.5      29      29.07      29.3
## 0.001219512 0.001219512 0.001219512 0.001219512 0.001219512 0.001219512
##      30      31      32      32.2      32.9      33
## 0.002439024 0.001219512 0.001219512 0.001219512 0.001219512 0.001219512
##      33.33      35      35.5      35.7      37.5      38
## 0.001219512 0.001219512 0.001219512 0.001219512 0.001219512 0.001219512
##      39      39.3      39.33      39.5      39.75      40
## 0.001219512 0.001219512 0.001219512 0.001219512 0.001219512 0.002439024
##      40.5      42.7      45      47      48      49.15
## 0.001219512 0.001219512 0.046341463 0.001219512 0.002439024 0.002439024
##      49.5      49.57      49.75      49.8      50      50.5
## 0.001219512 0.017073171 0.001219512 0.006097561 0.007317073 0.002439024
##      51      52      53      53.64      54      54.5
## 0.002439024 0.021951220 0.001219512 0.001219512 0.001219512 0.002439024
```

```

##          55          56.8          57.33          58          58.5          58.57
## 0.001219512 0.003658537 0.010975610 0.002439024 0.001219512 0.001219512
##          59.5          60
## 0.001219512 0.003658537
##
## Conditional probabilities:
##          Distance_In_Km
## Y          [,1]          [,2]
## 0.01      2.5479009          NA
## 0.11      2.6177784          NA
## 2.5      22.0808761 22.912853
## 2.9      37.7251378 27.070153
## 3        34.0447226 26.426231
## 3.25      0.4189497          NA
## 3.3      33.0573410 27.472687
## 3.4      13.8409749          NA
## 3.5      28.7409726 21.188613
## 3.7      25.8192337 26.751145
## 3.8      39.0022651          NA
## 4         39.2228944 32.357884
## 4.1      25.0725978 22.363298
## 4.5      20.7843395 22.275927
## 4.9      20.3565608 26.408721
## 5        17.3290563 22.180390
## 5.1       0.7910048          NA
## 5.3      24.8641845 29.274058
## 5.5      17.6310155 25.550488
## 5.7      38.0005633 24.921268
## 6         5.8908103  6.706961
## 6.1       5.3841191  6.429328
## 6.2       0.1113195          NA
## 6.3      38.6222069          NA
## 6.5      35.3511988 27.392645
## 6.9      23.5237228 25.333211
## 7        12.6479180 13.987339
## 7.3      38.4285477 27.306757
## 7.5      31.6824242 33.862995
## 7.7      26.4652365 15.365485
## 8        19.4825927 18.605070
## 8.1      46.5446445 27.403782
## 8.2      10.9496570          NA
## 8.5      40.3814211 20.207567
## 8.6      53.0065837          NA
## 8.9      60.4118282 12.954369
## 9        25.7655882 18.173573
## 9.3      21.6479891 21.030324
## 9.5      36.8210795 39.082656
## 9.6       0.7792364          NA
## 9.7      43.8028491 15.087914
## 9.8       3.1958325          NA
## 10       46.9797054 30.284621
## 10.1     25.0079932 20.927732

```



##	10.2	0.6064151	NA
##	10.3	1.0103798	NA
##	10.5	33.5703407	31.169566
##	10.9	26.0425422	NA
##	11	52.8850691	26.644111
##	11.3	36.1270052	24.257916
##	11.5	41.5466553	NA
##	11.7	12.0194889	7.388253
##	12	44.1006464	36.819044
##	12.1	19.7886938	13.935128
##	12.5	45.5540423	29.524464
##	12.9	7.6497091	10.423294
##	13	37.4942128	24.596267
##	13.3	13.1846810	NA
##	13.5	17.7674502	13.380645
##	13.8	0.7972678	NA
##	14	27.8763817	25.804370
##	14.1	52.5440523	8.518843
##	14.5	29.1652861	21.022864
##	14.6	54.8227250	NA
##	14.9	19.6471716	20.595603
##	15	2.8195232	1.100715
##	15.5	12.1418029	15.703828
##	15.7	9.9401851	NA
##	15.9	20.1596833	NA
##	16	64.5070556	NA
##	16.5	30.9218245	23.033063
##	16.9	27.4040126	33.224435
##	17	39.2888415	27.736496
##	17.3	0.2226390	NA
##	17.5	27.9379725	25.262726
##	17.9	29.3443625	NA
##	18	4.9584936	NA
##	18.1	0.3373876	NA
##	18.9	43.3468102	NA
##	19	3.7079537	NA
##	19.3	56.2555520	NA
##	19.7	49.4254984	NA
##	19.8	0.4768149	NA
##	20	12.3279851	22.412380
##	20.5	28.0313821	37.240413
##	20.9	54.5514320	31.240035
##	23	34.7792858	NA
##	23.7	62.5384993	NA
##	24.5	4.2576416	NA
##	25	14.1361486	18.790441
##	27	41.4726316	NA
##	28	0.4770395	NA
##	28.5	42.8602912	NA
##	29	23.5022841	NA
##	29.07	61.6608020	NA
##	29.3	27.3202839	NA

```

## 30      14.2162677  7.302095
## 31      1.3453087      NA
## 32      0.1689370      NA
## 32.2    4.1331101      NA
## 32.9    48.2727208      NA
## 33      70.3464558      NA
## 33.33   0.3369768      NA
## 35      0.4215060      NA
## 35.5    4.0968217      NA
## 35.7    69.4029453      NA
## 37.5    16.2146774      NA
## 38      3.1966026      NA
## 39      32.7575902      NA
## 39.3    71.1297560      NA
## 39.33   27.1882054      NA
## 39.5    56.1052094      NA
## 39.75   51.5469315      NA
## 40      1.2047137  1.308768
## 40.5    22.2395348      NA
## 42.7    0.3444593      NA
## 45      18.2949010  21.766500
## 47      0.3370903      NA
## 48      3.1114200  3.828580
## 49.15   8.5623809  3.056398
## 49.5    9.6229690      NA
## 49.57   4.7925190  3.857396
## 49.75   0.6433941      NA
## 49.8    6.5965257  6.544769
## 50      1.4089737  2.318874
## 50.5    11.6526168  6.823416
## 51      3.3334793  1.130297
## 52      13.4893381  24.157310
## 53      7.8215156      NA
## 53.64   28.0296990      NA
## 54      5.2407085      NA
## 54.5    2.0926973  1.800871
## 55      0.5816101      NA
## 56.8    23.2305708  35.868620
## 57.33   2.2508692  2.738756
## 58      1.6710081  1.576015
## 58.5    0.4042437      NA
## 58.57   49.2121304      NA
## 59.5    6.8566415      NA
## 60      10.6152528  15.110639

```

This output has given us various a priori probabilities and indicates that the model has found around 60 classes, each corresponding to a different, tiny probability. Because of this, the training data may appear scattered within most classes for the corresponding variables.

## Prediction and Evaluation

First, let's predict and evaluate on the test data of the logistic regression model.

```
p1 <- predict(glm1, newdata=test, type="response")
pred <- ifelse(p1 > 0.5, 1, 0)
acc <- mean(p1==test$fare_amount)
print(paste("accuracy =", acc))
```

```
## [1] "accuracy = 0"
```

Now, let's predict and evaluate on the test data of the Naive Bayes model.

```
p2 <- predict(nbl, newdata=test, type="class")
acc2 <- mean(p2==test$fare_amount)
print(paste("accuracy =", acc2))
```

```
## [1] "accuracy = 0.344660194174757"
```

If we assume the logistic regression model had an accuracy of 1.0 compared to the 0.345 accuracies provided by the Naive Bayes model. Based on these accuracy numbers, the logistic regression model could correctly predict all test samples, especially since the mean of the `pred` variable was 1. The logistic regression model would be a better fit for the data. This might be because there is an approximately linear relationship between the features and target variable rather than independence between features assumed by the Naive Bayes model.

## Naive Bayes vs Logistic Regression

The strengths of Naive Bayes come from being useful for classification with a large dataset and independence between features. It works well with scattered data and can typically eliminate features that are not relevant to the prediction. However, its weaknesses fall under the fact that there might be a strong relationship between all the features, especially when there is not necessarily an independence between them. Logistic regression also finds its strengths in classification and a large dataset. It is easy to implement and visualize the impact of features on a target variable. However, there might only sometimes be a linear relationship among features. It only works well if there are a few outliers in the dataset so it may require further cleaning.

### Sources

- <https://dplyr.tidyverse.org/index.html> (<https://dplyr.tidyverse.org/index.html>)
- <https://www.rdocumentation.org/packages/geosphere/versions/1.5-18>  
(<https://www.rdocumentation.org/packages/geosphere/versions/1.5-18>)
- <https://r4ds.had.co.nz/transform.html#transform> (<https://r4ds.had.co.nz/transform.html#transform>)
- <https://stackoverflow.com/questions/32363998/function-to-calculate-geospatial-distance-between-two-points-lat-long-using-r> (<https://stackoverflow.com/questions/32363998/function-to-calculate-geospatial-distance-between-two-points-lat-long-using-r>)
- <https://stackoverflow.com/questions/40554592/geosphere-disthaversine-dplyr-error-wrong-length-for-vector-should-be-2> (<https://stackoverflow.com/questions/40554592/geosphere-disthaversine-dplyr-error-wrong-length-for-vector-should-be-2>)
- <https://www.youtube.com/watch?v=MHbzCs05Luo> (<https://www.youtube.com/watch?v=MHbzCs05Luo>)
- [https://github.com/kjmazidi/Machine\\_Learning\\_2nd\\_edition/tree/master/Part\\_2\\_Linear\\_Models](https://github.com/kjmazidi/Machine_Learning_2nd_edition/tree/master/Part_2_Linear_Models)  
([https://github.com/kjmazidi/Machine\\_Learning\\_2nd\\_edition/tree/master/Part\\_2\\_Linear\\_Models](https://github.com/kjmazidi/Machine_Learning_2nd_edition/tree/master/Part_2_Linear_Models))
- <http://www.sthda.com/english/articles/36-classification-methods-essentials/151-logistic-regression-essentials-in-r/> (<http://www.sthda.com/english/articles/36-classification-methods-essentials/151-logistic-regression-essentials-in-r/>)